

OPERATING SYSTEM

Enterprise Decision Readiness Report

A deep field report for leaders deciding what to fix first when mandate clarity, evidence trust, AI judgment, ownership, and organizational capability are all part of the same problem.



Generated editorial image · a decision room where evidence, people, and operational context occupy the same frame.

Format: executive field report **Use:** board, executive workshop, or portfolio review **Read:** 45–55 minutes

Working thesis: systems of decision are becoming the next operating layer

EXECUTIVE READOUT

Decision readiness is not the presence of a strategy, a dashboard, or a capable team in isolation.

It is the organization's ability to name the next consequential decision, show the evidence behind it, assign an accountable owner, and move through the required exception path without losing trust. When one element is missing, delay is often misdiagnosed as a technology, process, or talent problem. The first intervention should usually restore the connection between signal, judgment, action, and learning.

INTERPRETATION

This report's strong view is that the enterprise is moving from a world organized mainly around systems of record and systems of engagement toward a third layer: **systems of decision**

. These will not be one magical platform. They will be an operating arrangement in which governed evidence, human judgment, AI assistance, workflow controls, and feedback are deliberately connected. Organizations that build this layer well will make better decisions faster. Organizations that automate fragments without building the layer will create faster ambiguity.

How to use this report

Use the diagnostic with one consequential decision rather than scoring the whole enterprise at once. Choose a portfolio, service, modernization move, operating-model change, or AI-enabled workflow that is important enough to expose the real system. Invite the person accountable for the outcome, the people who produce or use the evidence, the technology or data owner, someone who understands the change burden on the front line, and—when the consequence warrants it—a risk, privacy, security, finance, or legal partner.

Run the five-part readout in one session, then schedule a second session to test the evidence and ownership claims against actual work. The goal is not a perfect score. The goal is a shared view of the first constraint, the owner of the next move, the decision boundary that must remain human, and the evidence that will tell the team whether the intervention is working.

There is a temptation to make readiness a broad maturity assessment. Resist that temptation at the beginning. A broad assessment produces a broad backlog. A specific decision produces an operating truth. Once one decision is made more visible, repeatable, and accountable, the organization can decide which patterns deserve to spread.

The thesis: the next enterprise advantage is decision infrastructure

For decades, enterprise investment has accumulated around systems of record: financial, customer, workforce, supply, asset, and case data. Then systems of engagement made those records easier to reach through portals, mobile experiences, collaboration tools, and APIs. The emerging question is different: *how does the organization turn a changing signal into a responsible decision and a measured action?*

That question is becoming urgent because the volume and velocity of signals are increasing while accountability is not becoming optional. AI can summarize, classify, forecast, recommend, generate, and sometimes execute. Automation can route work across boundaries. Data products can make evidence more reusable. But none of those capabilities answers who is allowed to decide, what evidence is sufficient, what uncertainty is acceptable, how an exception is handled, or who carries the consequence when the system is wrong.

The future-facing implication is direct: **the winning enterprise will not simply have more AI, more data, or more dashboards. It will have a better designed decision surface.** The decision surface is where a person encounters the mandate, the current evidence, the model's contribution, the known limits, the available choices, the approval boundary, and the next action. It may be a case screen, a planning meeting, a control room, a clinical workflow, a finance review, or an executive portfolio ritual. Its quality will increasingly determine whether technology creates leverage or just creates more things to supervise.

From project delivery to decision throughput

Interpretation: leaders will ask less often whether a system shipped and more often whether a critical decision now moves with less avoidable delay, less rework, and clearer accountability. Delivery remains important, but it becomes a means to improve a decision loop.

From data ownership to evidence products

Interpretation: useful evidence will be managed like a product with a purpose, owner, freshness expectation, quality signal, access boundary, and retirement rule. A catalog is necessary in many environments; it is not the same thing as evidence that is ready for a decision.

From AI policy to AI operating design

Interpretation: policy will increasingly be judged by whether it appears in the workflow as permissions, review gates, evaluation records, escalation paths, and recovery practices. A policy that cannot be enacted at the point of work is a statement of intent, not a control system.

From adoption campaigns to capability transfer


Interpretation: the durable unit of change will be a role performing a decision differently, with practice and feedback. Communications may create awareness; capability is what lets the new behavior survive pressure, turnover, and changing conditions.

The three bets leaders should make now

- ☐ **Invest in decision visibility before broad automation.** Map a few high-consequence decisions end to end. Make the evidence path, owner, exception route, and action visible before asking a model to accelerate it.

- ☐ **Design for uncertainty as a first-class condition.** A future-ready system should expose freshness, confidence, missing data, disagreement, and override—not hide them behind a polished answer.

- **Build reusable control patterns, not one-off approvals.** Human review, access boundaries, source provenance, evaluation, incident response, and retirement should be composable patterns that can travel across use cases without pretending that every context is identical.



Interpretation: a decision system is a loop, not a dashboard. The organization creates value when the loop can move, pause, explain itself, recover from error, and learn from outcomes.

► Accessible text equivalent

The five-part readiness model

PART	QUESTION	VISIBLE EVIDENCE	TYPICAL FAILURE SIGNAL	FUTURE IMPLICATION
Mandate	What outcome and decision are we responsible for?	A plain-language promise, horizon, and escalation rule.	Every meeting restates the ambition; no one can state the choice.	Decision systems will be configured around outcomes and boundaries, not vague transformation themes.
Evidence	What must be true for a responsible decision?	Named sources, owners, freshness, caveats, and a route to the point of work.	Teams debate whose number is correct instead of what to do next.	Evidence quality and provenance will become a management responsibility, not only a data-team concern.
Judgment	Where should automation help and where must a person decide?	Human review, exception, override, and recovery paths.	A pilot is celebrated before accountability is designed.	AI value will be measured by the quality of the human-machine boundary, not by model novelty.

PART	QUESTION	VISIBLE EVIDENCE	TYPICAL FAILURE SIGNAL	FUTURE IMPLICATION
Ownership	Who can change the system when the signal changes?	Decision rights, service ownership, and a review cadence.	Responsibility is distributed across functions with no final owner.	Every durable automated service will need an operating owner with authority after launch.
Capability	Can the people carrying the work use and improve the new way?	Role design, practice, training, feedback, and transition support.	Adoption is treated as communications after the technical launch.	Capability transfer will be part of product design because the workflow is part of the product.

1. Mandate: make the decision smaller and more consequential

A mandate is ready when a team can state the service or outcome it owns, the decision that is blocked, the time horizon, and the consequence of doing nothing. “Modernize the enterprise” is a direction. “Reduce the distance between a trusted operational signal and the weekly capacity decision without weakening auditability” is a working mandate.

Ask the room to write the mandate in one sentence, then name what is explicitly outside the first release. A bounded mandate gives architecture, finance, risk, and change leaders something to coordinate around. It also prevents a broad transformation story from becoming an excuse to postpone the first useful decision.

Test it in the future state: if a capable AI system were available tomorrow, would the mandate tell it what good means, what it must not do, and who can overrule it? If not, the organization is not ready to automate the decision because it has not yet defined the decision.

Evidence to request: the one-sentence mandate, the in-scope decision, the out-of-scope boundary, the time horizon, the consequence of delay, and the person who can change the mandate when conditions change.

2. Evidence: follow the signal to the decision

Data maturity is often described as a platform question. Decision readiness asks a more exact question: can the responsible person find the signal, understand its provenance and limits, and use it at the moment a choice has to be made? Public data estates demonstrate why this matters. Data.gov's large catalog is useful infrastructure, but a catalog alone does not make a dataset trusted, current, governed, or relevant to a service decision.

Trace one decision from source to action. Record where the data is created, transformed, approved, displayed, challenged, and retained. Include the human interpretation layer. If the organization cannot draw this path, the first investment may be ownership and evidence design rather than another tool.

Future-facing implication: evidence will increasingly need to travel with its context. A recommendation without source, timestamp, confidence, scope, and known exclusions will be expensive to trust. The enterprise will need evidence products that can answer not only "what is the number?" but also "what decision is this number fit for?"

That does not mean every user needs a data lineage graph on screen. It means the organization must be able to retrieve the relevant explanation when a decision is challenged. The visible experience can be simple; the underlying path must be real.

Evidence to request: a source map, data owner, freshness rule, quality signal, transformation record, access boundary, definition of "fit for use," and a named process for correcting or retiring the evidence.

3. Judgment: couple AI to the workflow it changes

AI readiness is not a count of experiments. It is the ability to place an intelligent capability inside a real workflow with a clear purpose, a measurable failure boundary, and a human role that is neither ceremonial nor overloaded. Start with the decision and the consequence. Then define what the system may recommend, what it may execute, what requires review, and what happens when the signal is missing, stale, or wrong.

For each use case, document the source of truth, evaluation method, expected users, override path, monitoring owner, and recovery action. A useful control does not merely prohibit risk; it helps the right person make a safer decision quickly. An approval button that a reviewer cannot meaningfully assess is not human oversight. It is an accountability transfer disguised as a control.

Interpretation: the near-term competitive advantage will come from workflow-native AI rather than free-floating assistants. The assistant that knows the decision boundary, sees the right evidence, cites its basis, and can hand an exception to a named owner will be more valuable than a general-purpose answer engine operating outside the work.

Counter-signal: not every workflow needs AI. A simpler rule, better data definition, clearer ownership, or redesigned queue may produce more value with less operational risk. The right question is not “where can we put a model?” It is “where does a bounded form of intelligence improve a decision enough to justify the new control surface?”

Evidence to request: use-case purpose, prohibited uses, evaluation set or test method, expected failure modes, human review design, permissions, logging, incident path, rollback or disablement procedure, and retirement trigger.

4. Ownership: turn accountability into a service

Ownership is more than a name on a RACI chart. A real owner has authority to change the service, access to the evidence, a budget or escalation path, and a cadence in which the performance of the system is reviewed. If the work crosses functions, the team should name one outcome owner and the supporting control owners around it.

Make the review rhythm explicit. A weekly operating review may examine exceptions and flow. A monthly portfolio review may examine value, risk, and dependency. A quarterly leadership review may decide whether to scale, pause, redesign, or retire. Each cadence needs a decision to make, not only a report to receive.

Future-facing implication: as more decisions become partially automated, post-launch ownership will become a differentiator. The implementation team may leave, the model may change, the data may drift, and the policy may evolve. A service owner must remain able to

inspect, adjust, and stop the system. “The vendor owns the model” cannot be the organization’s operating model.

Evidence to request: outcome owner, control owners, decision rights, escalation chain, service-level expectations, review calendar, change authority, budget route, and a named person who can pause the capability.

5. Capability: make the new behavior easier to carry

Capability is the bridge between a sound design and a durable result. Map the roles that experience the changed workflow, the decisions they must make differently, the evidence they need, and the feedback they can return to the service owner. Training is one input. Practice, coaching, role clarity, and a safe path to challenge the system are the operating capability.

Capability gaps often appear before cost or compliance gaps. Watch for manual workarounds, duplicate review, unowned exceptions, low-quality inputs, and local spreadsheets that quietly become the real system. Those signals show where the design has not yet met the work.

Interpretation: the organizations that benefit from AI will be the ones that treat judgment as a craft to be strengthened, not a human bottleneck to be eliminated. People will need to know when to trust a recommendation, when to investigate, how to record an exception, and how to improve the source. The future role is not “human in the loop” as a slogan; it is a person with enough context, time, authority, and practice to exercise judgment.

Evidence to request: role map, changed decisions, practice environment, enablement plan, feedback channel, accommodations, frontline operating constraints, and measures of whether the new way is actually easier and safer to carry.

Where enterprise decision systems are heading

The following horizon is an interpretation, not a forecast with a guaranteed timetable. It is useful because it gives leaders a direction for designing today's work. The key is to build reversible foundations: clear mandates, inspectable evidence, explicit ownership, and bounded experiments. Those foundations remain useful even if a particular model, vendor, or market assumption changes.

HORIZON	WHAT BECOMES NORMAL	WHAT LEADERS SHOULD BUILD NOW	WHAT TO WATCH
Near term	Assistive systems summarize context, prepare options, find exceptions, and route work across existing teams.	Decision records, source visibility, evaluation discipline, clear review boundaries, and an owner after launch.	People accept fluent outputs without checking provenance; pilots multiply without shared controls.
Next phase	Workflow-native agents coordinate bounded tasks across systems with permissions and explicit handoffs.	Service maps, machine-readable policies, identity and access boundaries, event logs, recovery paths, and operating cadences.	Local automations become invisible infrastructure; exception volume exceeds human capacity.
Longer horizon	Organizations compose decision services that continuously compare signals, options, outcomes, and policy constraints.	A culture of evidence correction, model and rule retirement, portfolio-level risk review, and capability transfer.	Optimization narrows the field of view; speed outpaces legitimacy; ownership fragments across vendors and functions.

The practical lesson is not to wait for a fully autonomous enterprise. It is to design for the direction of travel while protecting reversibility. Every new capability should have a clear purpose, a bounded authority, an observable output, a human escalation path, and a way to turn it off. That is how an organization can move quickly without confusing novelty with readiness.

Industry implications

Decision readiness is portable, but the consequence of error is not. The same operating pattern will look different in a public service, a hospital, a factory, a financial institution, or a professional-services firm. The categories below are directional interpretations to help a leadership team ask better questions. They are not claims about any particular organization or a substitute for sector-specific obligations.

Regulated and public-facing services

Interpretation: the differentiator will be explainable service decisions that preserve dignity and recourse while reducing administrative drag. Automation should help staff see the relevant context and route work; it should not make the organization unable to explain or challenge an outcome. Readiness therefore includes an appeal path, a clear record of the evidence used, and an owner who can correct systemic patterns rather than only individual cases.

Leadership question: if the affected person, an ombudsperson, an auditor, or a frontline worker challenges the decision, can the organization reconstruct the path without relying on the memory of one operator?

Healthcare and human services

Interpretation: the highest-value systems will reduce cognitive and administrative burden while keeping professional judgment visible. The workflow needs to make uncertainty and disagreement easy to surface. A recommendation that saves time but increases reconciliation, duplicate documentation, or hidden review work is not a successful decision system.

Leadership question: does the design give the professional enough context and enough time to disagree safely, and does the organization learn from those disagreements?

Industrial, energy, and asset-intensive operations

Interpretation: decision readiness will be expressed through the connection between operational signals, maintenance or capacity choices, safety boundaries, and field action. The value is not only predictive accuracy. It is whether the right person receives a usable recommendation before the operating window closes, with a safe fallback when sensors, networks, or assumptions fail.

Leadership question: what happens when the data is unavailable in the exact moment the system is expected to help, and who is trained and authorized to switch to the fallback?

Financial and risk-sensitive enterprises

Interpretation: decision systems will increasingly be evaluated as portfolios of policy, evidence, model behavior, and human override. The strongest institutions will distinguish a decision that may be automated from one that may merely be prepared or prioritized. They will also make it possible to inspect drift across segments, time periods, and operating conditions rather than celebrating a single average measure.

Leadership question: can the organization tell the difference between a model performing badly, a policy changing, a data source drifting, and a human process quietly compensating for all three?

Professional services and knowledge work

Interpretation: the opportunity is to move senior judgment toward the questions where it has the highest value while making routine preparation more consistent. The risk is that generated work appears finished before it is sufficiently tested. Firms will need provenance, review rituals, client-specific boundaries, and a way to teach newer practitioners how to interrogate a result rather than merely edit its prose.

Leadership question: does automation increase the organization's capacity to develop judgment, or does it hide the learning path and concentrate quality in a few reviewers?

Global and distributed operating models

Interpretation: decision systems can reduce the friction of time zones, handoffs, and fragmented evidence, but only if they make local context visible. A globally standardized signal may be useful for comparison and wrong for action. The future operating model will likely combine shared definitions and controls with local authority to interpret conditions and raise exceptions.

Leadership question: which parts of the decision must be consistent everywhere, and where must local teams retain the right to disagree with the global signal?

How to prioritize the portfolio

Most organizations do not lack potential use cases. They lack a disciplined way to decide which ones deserve attention first. A useful portfolio screen considers consequence, decision frequency, evidence availability, reversibility, control complexity, and capability transfer. Do not rank only by estimated efficiency. A small improvement in a frequent decision may matter more than a large improvement in a rare decision; a reversible workflow may be a better proving ground than a spectacular but irreversible one.

PORTFOLIO LENS	PROMPT	PREFER EARLY WORK THAT...	PAUSE WHEN...
Consequence	What happens if the decision is wrong or delayed?	Has meaningful value and a clearly bounded harm surface.	The team cannot name who bears the consequence.
Frequency	How often does the decision occur?	Creates repeated learning and visible operating feedback.	It is so rare that evidence will not accumulate.
Evidence	Can the signal be traced, challenged, and corrected?	Has a known owner and enough quality to support a bounded test.	Data quality is being treated as a future clean-up task.
Reversibility	Can the organization stop or unwind the intervention?	Has a fallback and a clear stop signal.	Failure would be durable before learning is possible.
Capability	Will the people doing the work be able to improve it?	Includes practice, feedback, and local champions with authority.	Adoption depends on silent workarounds or heroic individuals.

Use the screen to choose a sequence, not to produce a false numerical precision. The first candidate should usually be consequential enough to matter, bounded enough to learn from, and visible enough that leaders can inspect the whole loop. The second candidate should test a different constraint. A portfolio becomes strategic when it teaches the organization how to improve its decision infrastructure, not only when it accumulates successful pilots.

Measures that reveal readiness

Readiness measures should expose the health of the decision loop. Avoid measuring only activity—number of pilots, dashboards, prompts, or trained users. Those can be useful leading indicators, but they do not prove that a decision improved. Define measures before launch and include at least one signal from each of four perspectives: flow, quality, control, and capability.

☐ **Flow:** time from signal to decision, handoff count, avoidable queue time, or the proportion of cases that reach the right owner without rework.

☐ **Quality:** correction rate, disagreement rate, completeness of the decision record, outcome consistency, or the percentage of recommendations supported by usable evidence.

☐ **Control:** exception aging, override review, access violations, unresolved incidents, or the time required to disable or roll back the capability.

☐ **Capability:** successful practice, confidence to challenge the system, adoption of the intended workflow, quality of feedback, or reduction in shadow processes.

Interpretation: disagreement is not automatically a failure. A healthy decision system may increase visible disagreement because people can finally see the basis of a recommendation and challenge it. The more useful measure is whether disagreement is resolved with evidence, recorded when material, and used to improve the system.

Define a small number of measures that a real owner can review and act upon. If a metric cannot trigger a decision, it is probably an observation rather than an operating measure.

The executive workshop

The following workshop is designed for 90 to 120 minutes. It is deliberately demanding: the aim is to leave with a decision boundary and a first move, not a shared feeling that the topic is important.

TIME	ACTIVITY	OUTPUT	FACILITATOR PRESSURE-TEST
10 min	Frame the consequential decision.	One sentence naming outcome, decision, horizon, and consequence of delay.	What choice will be different if this work succeeds?
20 min	Trace evidence to action.	Source-to-decision map with gaps and caveats visible.	Where does interpretation enter, and who can challenge it?
20 min	Draw the judgment boundary.	Allowed recommendation, allowed execution, mandatory review, fallback.	What is the worst plausible failure and how is it detected?
15 min	Name ownership and cadence.	Outcome owner, control owners, escalation, review decision.	Who can pause the system on a difficult day?
15 min	Map capability transfer.	Changed roles, practice needs, feedback path, shadow work.	What becomes easier for the person carrying the work?
10–30 min	Choose the bounded test.	First slice, measures, stop/continue signal, next meeting.	What will we refuse to automate until we learn more?

Workshop roles

Ask the outcome owner to open and close the session. Ask an evidence owner to bring the actual source path, not a presentation about the source path. Ask a frontline or service representative to describe the work as performed, including workarounds. Ask the technology or AI lead to state what the system can and cannot do. Ask a control partner to identify the consequence that would change the review boundary. These roles can be held by different people or combined in a small team, but the perspectives should be present.

Pre-work

☐ Bring one recent decision record, case, or portfolio choice from the workflow.

☐ Bring the current source definitions, freshness expectations, and known quality limitations.

☐ Bring the current process or service map, including manual workarounds and escalations.

☐ Bring any existing evaluation, incident, privacy, security, or review record relevant to the proposed change.

☐ Ask participants to write what they believe the first decision should be before the meeting; compare the answers at the start.

Decision readiness worksheet

1. The decision we are making:

2. The outcome and consequence of delay:

3. The evidence we trust, including freshness and limits:

4. What the system may recommend:

5. What the system may execute:

6. What always requires human review:

7. The fallback when evidence, model, or workflow fails:

8. Outcome owner and control owners:

9. First bounded test:

10. Stop, continue, or redesign signal:

11. Capability that must transfer to the people doing the work:

--

Counterpoints: what could change the conclusion?

A strong thesis is useful only if it can be challenged. The conclusion that decision infrastructure will become a major enterprise advantage could change, narrow, or arrive differently under several conditions:

- ☐ **Economics may favor simplification.** If the cost and complexity of advanced AI remain higher than the value of the decisions it improves, many organizations will rationally choose clearer rules, better workflow design, or conventional analytics. That would not invalidate decision readiness; it would reinforce the principle that the operating loop matters more than the tool.
- ☐ **Trust may become the binding constraint.** A high-profile failure, a security incident, or an unacceptable pattern of unexplained outcomes could slow deployment. Organizations with inspectable evidence, meaningful recourse, and the ability to pause will be better positioned, but speed will not always be the responsible goal.
- ☐ **Regulatory and contractual requirements may diverge.** Different jurisdictions, sectors, customers, and labor arrangements may require different levels of explanation, human review, retention, or localization. A universal template will fail; portable principles with local control design will be more resilient.
- ☐ **Data access may remain structurally limited.** Some valuable decisions are constrained by sparse, delayed, proprietary, or highly contextual evidence. A model cannot recover information the operating model does not collect. The better response may be to change the service, measurement, or relationship—not to add inference.

- **People may reject the workflow for good reasons.** Low adoption may signal poor enablement, but it may also reveal that the new way removes necessary judgment, shifts risk unfairly, or makes work harder. A readiness practice must treat frontline disagreement as evidence, not noise.

These counterpoints suggest a practical rule: **build decision infrastructure as a reversible capability, not as an irreversible bet on a technology category.** Preserve the source map, ownership, review boundary, fallback, and learning cadence even if the chosen model, vendor, or automation approach changes.

A first 90-day sequence

The sequence below is intentionally modest. It is designed to produce operating evidence rather than a large transformation promise.

Days 1–15: choose and frame

Select one decision with a real consequence and a visible owner. Write the mandate, decision boundary, out-of-scope list, and consequence of delay. Collect one recent real example. If the group cannot agree on the decision, stop and resolve that before designing technology.

Days 16–30: trace and expose

Map the evidence from creation to action. Name the source owner, freshness expectation, transformation, interpretation, access boundary, and correction path. Draw the current workflow including exceptions and shadow work. Record what is unknown. Unknowns are not failures; hidden unknowns are.

Days 31–60: test the boundary

Run a bounded test with a narrow user group, service slice, or decision interval. Test the human-machine boundary more aggressively than the happy path. Ask participants to challenge the recommendation, operate with missing evidence, and use the fallback. Measure flow, quality, control, and capability. Do not scale because the demonstration was impressive.

Days 61–90: decide what deserves to spread

Review the results with the outcome owner and control owners. Decide whether to continue, redesign, pause, or retire. Extract reusable patterns: evidence requirements, review gates, logging, permissions, exception handling, training, and the cadence. Put those patterns into the

normal management rhythm. If the organization cannot name the next owner after the project team leaves, it is not ready to scale.

Evidence and source notes

SOURCE-GROUNDED CONTEXT

The links below anchor the report's discussion of public goals, data infrastructure, and responsible AI risk management. They support the questions and framing; they do not validate a particular organization's maturity, controls, or outcomes.

☐ United Nations · Sustainable Development Goals Report 2026 — source context for the report's emphasis on measurable outcomes, institutional capacity, and the distance between public ambition and implementation.

☐ Data.gov · public data catalog — source context for the distinction between discoverable data infrastructure and evidence that is actually fit for a specific decision.

☐ NIST · AI Risk Management Framework — source context for treating AI risk as a lifecycle and governance concern rather than a one-time technical review.

☐ World Economic Forum · Future of Jobs Report 2025 — context for treating capability, role redesign, and organizational learning as part of readiness rather than a downstream training task.

Interpretation in this report includes the systems-of-decision thesis, the horizon, industry implications, portfolio recommendations, suggested measures, and the 90-day sequence. Those passages are strategic judgment intended to provoke a useful leadership conversation. They should be tested against the organization's own evidence, constraints, and obligations before adoption.

Source-grounded language is used narrowly. A linked source provides context, not a universal prescription. The report does not claim that any source endorses Global Enterprise's model, predicts a specific market outcome, or proves a customer result.

Limits and interpretation

This report is a decision aid, not a maturity certification, audit opinion, legal recommendation, security assessment, financial opinion, clinical protocol, or substitute for domain-specific privacy, regulatory, safety, or procurement advice. The model is intentionally portable. The evidence, risk boundary, decision rights, and measures must be adapted to the institution and the consequence of the work.

The report uses future-facing language because leaders need a direction for investment, but no horizon is guaranteed. Technology capability, economics, regulation, public trust, workforce practice, and geopolitical conditions can change the conclusion. Build foundations that remain valuable if the specific forecast is wrong: clear mandates, inspectable evidence, accountable ownership, reversible interventions, meaningful recourse, and the ability to learn.

Public sources can change. Verify source status, applicability, and current obligations before relying on any statement for a material decision. Do not place confidential, regulated, personal, or security-sensitive information into an unapproved public channel to complete this worksheet.

PUBLIC RESEARCH BASIS

- United Nations · Sustainable Development Goals Report 2026
- Data.gov · public data catalog
- NIST · AI Risk Management Framework
- World Economic Forum · Future of Jobs Report 2025

CONTINUE THE CONVERSATION

Bring the decision into the room.

Global Enterprise helps leaders make complex systems more legible, governable, and ready to move.

Request a principal dialogue · info@globalenterprise.com · +1 929 229 2918

